

Miscellaneous

IEEE 754 Floating Point Special Values

Special values defined in numpy: nan, inf,

NaNs can be used as a poor-man’s mask (if you don’t care what the original value was)

Note: cannot use equality to test NaNs. E.g.:

```
>>> myarr = np.array([1., 0., np.nan, 3.])
>>> np.nonzero(myarr == np.nan)
(array([], dtype=int64),)
>>> np.nan == np.nan # is always False! Use special numpy functions instead.
False
>>> myarr[myarr == np.nan] = 0. # doesn't work
>>> myarr
array([ 1.,  0., nan,  3.])
>>> myarr[np.isnan(myarr)] = 0. # use this instead find
>>> myarr
array([1.,  0.,  0.,  3.])
```

Other related special value functions:

```
isinf(): True if value is inf
isfinite(): True if not nan or inf
nan_to_num(): Map nan to 0, inf to max float, -inf to min float
```

The following corresponds to the usual functions except that nans are excluded from the results:

```
nansum()
nanmax()
nanmin()
nanargmax()
nanargmin()

>>> x = np.arange(10.)
>>> x[3] = np.nan
>>> x.sum()
nan
>>> np.nansum(x)
42.0
```

How numpy handles numerical exceptions

The default is to 'warn' for invalid, divide, and overflow and 'ignore' for underflow. But this can be changed, and it can be set individually for different kinds of exceptions. The different behaviors are:

- ‘ignore’: Take no action when the exception occurs.
- ‘warn’: Print a *RuntimeWarning* (via the Python [warnings](#) module).
- ‘raise’: Raise a *FloatingPointError*.
- ‘call’: Call a function specified using the *seterrcall* function.
- ‘print’: Print a warning directly to *stdout*.
- ‘log’: Record error in a Log object specified by *seterrcall*.

🔍 Search the docs ...

- [What is NumPy?](#)
- [Installation](#)
- [NumPy quickstart](#)
- [NumPy: the absolute basics for beginners](#)
- [NumPy fundamentals](#)
- [Miscellaneous](#)**
- [NumPy for MATLAB users](#)
- [Building from source](#)
- [Using NumPy C-API](#)
- [NumPy Tutorials](#)
- [NumPy How Tos](#)
- [For downstream package authors](#)
- [F2PY user guide and reference manual](#)
- [Glossary](#)
- [Under-the-hood Documentation for developers](#)
- [Reporting bugs](#)
- [Release notes](#)
- [NumPy license](#)

These behaviors can be set for all kinds of errors or specific ones:

- all : apply to all numeric exceptions
- invalid : when NaNs are generated
- divide : divide by zero (for integers as well!)
- overflow : floating point overflows
- underflow : floating point underflows

Note that integer divide-by-zero is handled by the same machinery. These behaviors are set on a per-thread basis.

Examples

```
>>> oldsettings = np.seterr(all='warn')
>>> np.zeros(5,dtype=np.float32)/0.
Traceback (most recent call last):
...
RuntimeWarning: invalid value encountered in divide
>>> j = np.seterr(under='ignore')
>>> np.array([1.e-100])**10
array([0.])
>>> j = np.seterr(invalid='raise')
>>> np.sqrt(np.array([-1.]))
Traceback (most recent call last):
...
FloatingPointError: invalid value encountered in sqrt
>>> def errorhandler(errstr, errflag):
...     print("saw stupid error!")
>>> np.seterrcall(errorhandler)
>>> j = np.seterr(all='call')
>>> np.zeros(5, dtype=np.int32)/0
saw stupid error!
array([nan, nan, nan, nan, nan])
>>> j = np.seterr(**oldsettings) # restore previous
...                               # error-handling settings
```

Interfacing to C

Only a survey of the choices. Little detail on how each works.

1. Bare metal, wrap your own C-code manually.

- Plusses:
 - Efficient
 - No dependencies on other tools
- Minuses:
 - Lots of learning overhead:
 - need to learn basics of Python C API
 - need to learn basics of numpy C API
 - need to learn how to handle reference counting and love it.
 - Reference counting often difficult to get right.
 - getting it wrong leads to memory leaks, and worse, segfaults

2. Cython

- Plusses:
 - avoid learning C API's
 - no dealing with reference counting
 - can code in pseudo python and generate C code
 - can also interface to existing C code
 - should shield you from changes to Python C api
 - has become the de-facto standard within the scientific Python community
 - fast indexing support for arrays
- Minuses:
 - Can write code in non-standard form which may become obsolete
 - Not as flexible as manual wrapping

3. ctypes

- Plusses:
 - part of Python standard library
 - good for interfacing to existing shareable libraries, particularly Windows DLLs
 - avoids API/reference counting issues
 - good numpy support: arrays have all these in their ctypes attribute:

```
a.ctypes.data
a.ctypes.data_as
a.ctypes.shape
a.ctypes.shape_as
a.ctypes.strides
a.ctypes.strides_as
```

- Minuses:
 - can't use for writing code to be turned into C extensions, only a wrapper tool.

4. SWIG (automatic wrapper generator)

- Plusses:
 - around a long time
 - multiple scripting language support
 - C++ support
 - Good for wrapping large (many functions) existing C libraries
- Minuses:
 - generates lots of code between Python and the C code
 - can cause performance problems that are nearly impossible to optimize out
 - interface files can be hard to write
 - doesn't necessarily avoid reference counting issues or needing to know API's

5. Psyco

- Plusses:
 - Turns pure python into efficient machine code through jit-like optimizations
 - very fast when it optimizes well
- Minuses:
 - Only on intel (windows?)
 - Doesn't do much for numpy?

Interfacing to Fortran:

The clear choice to wrap Fortran code is [f2py](#).

Pyfort is an older alternative, but not supported any longer. Fwrap is a newer project that looked promising but isn't being developed any longer.

Interfacing to C++:

1. Cython
 2. CXX
 3. Boost.python
 4. SWIG
 5. SIP (used mainly in PyQt)

Previous

Interoperability with NumPy

Next

NumPy for MATLAB users