

🔍 Search the docs ...

[Installation](#)

[Package overview](#)

[Getting started tutorials](#) ^

[What kind of data does pandas handle?](#)

[How do I read and write tabular data?](#)

[How do I select a subset of a DataFrame ?](#)

[How to create plots in pandas?](#)

[How to create new columns derived from existing columns?](#)

[How to calculate summary statistics?](#)

[How to reshape the layout of tables?](#)

[How to combine data from multiple tables?](#)

[How to handle time series data with ease?](#)

[How to manipulate textual data?](#)

[Comparison with other tools](#) v

[Community tutorials](#)

In [1]: `import pandas as pd`

Data used for this tutorial:

Air quality Nitrate data

For this tutorial, air quality data about NO_2 is used, made available by [openaq](#) and downloaded using the [py-openaq](#) package.

The `air_quality_no2_long.csv` data set provides NO_2 values for the measurement stations *FR04014*, *BETR801* and *London Westminster* in respectively Paris, Antwerp and London.

To raw data

```
In [2]: air_quality_no2 = pd.read_csv("data/air_quality_no2_long.csv",
...:                                parse_dates=True)
...:
In [3]: air_quality_no2 = air_quality_no2[["date.utc", "location",
...:                                       "parameter", "value"]]
...:
In [4]: air_quality_no2.head()
Out[4]:
```

		date.utc	location	parameter	value
0	2019-06-21 00:00:00+00:00	FR04014	no2	20.0	
1	2019-06-20 23:00:00+00:00	FR04014	no2	21.8	
2	2019-06-20 22:00:00+00:00	FR04014	no2	26.5	
3	2019-06-20 21:00:00+00:00	FR04014	no2	24.9	
4	2019-06-20 20:00:00+00:00	FR04014	no2	21.4	

Air quality Particulate matter data

For this tutorial, air quality data about Particulate matter less than 2.5 micrometers is used, made available by [openaq](#) and downloaded using the [py-openaq](#) package.

The `air_quality_pm25_long.csv` data set provides PM_{25} values for the measurement stations *FR04014*, *BETR801* and *London Westminster* in respectively Paris, Antwerp and London.


To raw data

```
In [5]: air_quality_pm25 = pd.read_csv("data/air_quality_pm25_long.csv",
...:                                   parse_dates=True)
...:
In [6]: air_quality_pm25 = air_quality_pm25[["date.utc", "location",
...:                                       "parameter", "value"]]
...:
In [7]: air_quality_pm25.head()
Out[7]:
```

		date.utc	location	parameter	value
0	2019-06-18 06:00:00+00:00	BETR801	pm25	18.0	
1	2019-06-17 08:00:00+00:00	BETR801	pm25	6.5	
2	2019-06-17 07:00:00+00:00	BETR801	pm25	18.5	
3	2019-06-17 06:00:00+00:00	BETR801	pm25	16.0	
4	2019-06-17 05:00:00+00:00	BETR801	pm25	7.5	

How to combine data from multiple tables?

Concatenating objects



? I want to combine the measurements of NO_2 and PM_{25} , two tables with a similar structure, in a single table

```
In [8]: air_quality = pd.concat([air_quality_pm25, air_quality_no2], axis=0)

In [9]: air_quality.head()
Out[9]:
```

	date.utc	location	parameter	value
0	2019-06-18 06:00:00+00:00	BETR801	pm25	18.0
1	2019-06-17 08:00:00+00:00	BETR801	pm25	6.5
2	2019-06-17 07:00:00+00:00	BETR801	pm25	18.5
3	2019-06-17 06:00:00+00:00	BETR801	pm25	16.0
4	2019-06-17 05:00:00+00:00	BETR801	pm25	7.5

The `concat()` function performs concatenation operations of multiple tables along one of the axis (row-wise or column-wise).

By default concatenation is along axis 0, so the resulting table combines the rows of the input tables. Let's check the shape of the original and the concatenated tables to verify the operation:

```
In [10]: print('Shape of the `air_quality_pm25` table: ', air_quality_pm25.shape)
Shape of the `air_quality_pm25` table: (1110, 4)

In [11]: print('Shape of the `air_quality_no2` table: ', air_quality_no2.shape)
Shape of the `air_quality_no2` table: (2068, 4)

In [12]: print('Shape of the resulting `air_quality` table: ', air_quality.shape)
Shape of the resulting `air_quality` table: (3178, 4)
```

Hence, the resulting table has $3178 = 1110 + 2068$ rows.

Note

The **axis** argument will return in a number of pandas methods that can be applied **along an axis**. A **DataFrame** has two corresponding axes: the first running vertically downwards across rows (axis 0), and the second running horizontally across columns (axis 1). Most operations like concatenation or summary statistics are by default across rows (axis 0), but can be applied across columns as well.

Sorting the table on the datetime information illustrates also the combination of both tables, with the `parameter` column defining the origin of the table (either `no2` from table `air_quality_no2` or `pm25` from table `air_quality_pm25`):

```
In [13]: air_quality = air_quality.sort_values("date.utc")

In [14]: air_quality.head()
Out[14]:
```

	date.utc	location	parameter	value
2067	2019-05-07 01:00:00+00:00	London Westminster	no2	23.0
1003	2019-05-07 01:00:00+00:00	FR04014	no2	25.0
100	2019-05-07 01:00:00+00:00	BETR801	pm25	12.5
1098	2019-05-07 01:00:00+00:00	BETR801	no2	50.5
1109	2019-05-07 01:00:00+00:00	London Westminster	pm25	8.0

In this specific example, the `parameter` column provided by the data ensures that each of the original tables can be identified. This is not always the case. the `concat` function provides a convenient solution with the `keys` argument, adding an additional (hierarchical) row index. For example:

```
In [15]: air_quality_ = pd.concat([air_quality_pm25, air_quality_no2], keys=["PM25", "NO2"])
```

```
In [16]: air_quality_.head()
Out[16]:
```

			date.utc	location	parameter	value
PM25	0	2019-06-18	06:00:00+00:00	BETR801	pm25	18.0
	1	2019-06-17	08:00:00+00:00	BETR801	pm25	6.5
	2	2019-06-17	07:00:00+00:00	BETR801	pm25	18.5
	3	2019-06-17	06:00:00+00:00	BETR801	pm25	16.0
	4	2019-06-17	05:00:00+00:00	BETR801	pm25	7.5

Note

The existence of multiple row/column indices at the same time has not been mentioned within these tutorials. *Hierarchical indexing* or *MultiIndex* is an advanced and powerful pandas feature to analyze higher dimensional data.

Multi-indexing is out of scope for this pandas introduction. For the moment, remember that the function `reset_index` can be used to convert any level of an index to a column, e.g.

```
air_quality.reset_index(level=0)
```

To user guide

Feel free to dive into the world of multi-indexing at the user guide section on [advanced indexing](#).

To user guide

More options on table concatenation (row and column wise) and how `concat` can be used to define the logic (union or intersection) of the indexes on the other axes is provided at the section on [object concatenation](#).

Join tables using a common identifier



? Add the station coordinates, provided by the stations metadata table, to the corresponding rows in the measurements table.

Warning

The air quality measurement station coordinates are stored in a data file `air_quality_stations.csv`, downloaded using the [py-openaq](#) package.

```
In [17]: stations_coord = pd.read_csv("data/air_quality_stations.csv")

In [18]: stations_coord.head()
Out[18]:
```

	location	coordinates.latitude	coordinates.longitude
0	BELAL01	51.23619	4.38522
1	BELHB23	51.17030	4.34100
2	BELLD01	51.10998	5.00486
3	BELLD02	51.12038	5.02155
4	BELR833	51.32766	4.36226

Note

The stations used in this example (FR04014, BETR801 and London Westminster) are just three entries enlisted in the metadata table. We only want to add the coordinates of these three to the measurements table, each on the corresponding rows of the `air_quality` table.

```
In [19]: air_quality.head()
Out[19]:
```

		date.utc	location	parameter	value
2067	2019-05-07 01:00:00+00:00	London Westminster	no2	23.0	
1003	2019-05-07 01:00:00+00:00	FR04014	no2	25.0	
100	2019-05-07 01:00:00+00:00	BETR801	pm25	12.5	
1098	2019-05-07 01:00:00+00:00	BETR801	no2	50.5	
1109	2019-05-07 01:00:00+00:00	London Westminster	pm25	8.0	


```
In [20]: air_quality = pd.merge(air_quality, stations_coord, how="left", on="location")

In [21]: air_quality.head()
Out[21]:
```

		date.utc	location	parameter	value	coordinates.latitude
						coordinates.longitude
0	2019-05-07 01:00:00+00:00	London Westminster	no2	23.0	51.49467	-0.13193
1	2019-05-07 01:00:00+00:00	FR04014	no2	25.0	48.83724	2.39390
2	2019-05-07 01:00:00+00:00	FR04014	no2	25.0	48.83722	2.39390
3	2019-05-07 01:00:00+00:00	BETR801	pm25	12.5	51.20966	4.43182
4	2019-05-07 01:00:00+00:00	BETR801	no2	50.5	51.20966	4.43182

Using the `merge()` function, for each of the rows in the `air_quality` table, the corresponding coordinates are added from the `air_quality_stations_coord` table. Both tables have the column `location` in common which is used as a key to combine the information. By choosing the `left` join, only the locations available in the `air_quality` (left) table, i.e. FR04014, BETR801 and London Westminster, end up in the resulting table. The `merge` function supports multiple join options similar to database-style operations.

? Add the parameter full description and name, provided by the parameters metadata table, to the measurements table

 **Warning**

The air quality parameters metadata are stored in a data file `air_quality_parameters.csv`, downloaded using the `py-openaq` package.

```
In [22]: air_quality_parameters = pd.read_csv("data/air_quality_parameters.csv")

In [23]: air_quality_parameters.head()
Out[23]:
```

	id		description	name
0	bc		Black Carbon	BC
1	co		Carbon Monoxide	CO
2	no2		Nitrogen Dioxide	NO2
3	o3		Ozone	O3
4	pm10	Particulate matter less than 10 micrometers in...		PM10

```
In [24]: air_quality = pd.merge(air_quality, air_quality_parameters,
.....:                          how='left', left_on='parameter', right_on='id')
.....:

In [25]: air_quality.head()
Out[25]:
```

		date.utc	location	parameter	...	id	
						description	name
0	2019-05-07 01:00:00+00:00	London Westminster	no2	...	no2	Nitrogen Dioxide	NO2
1	2019-05-07 01:00:00+00:00	FR04014	no2	...	no2	Nitrogen Dioxide	NO2
2	2019-05-07 01:00:00+00:00	FR04014	no2	...	no2	Nitrogen Dioxide	NO2
3	2019-05-07 01:00:00+00:00	BETR801	pm25	...	pm25	Particulate matter less than 2.5 micrometers i...	PM2.5
4	2019-05-07 01:00:00+00:00	BETR801	no2	...	no2	Nitrogen Dioxide	NO2

[5 rows x 9 columns]

Compared to the previous example, there is no common column name. However, the `parameter` column in the `air_quality` table and the `id` column in the `air_quality_parameters_name` both provide the measured variable in a common format. The `left_on` and `right_on` arguments are used here (instead of just `on`) to make the link between the two tables.

pandas supports also inner, outer, and right joins. More information on join/merge of tables is provided in

To user guide

the user guide section on [database style merging of tables](#). Or have a look at the [comparison with SQL](#) page.

REMEMBER

- Multiple tables can be concatenated both column-wise and row-wise using the `concat` function.
- For database-like merging/joining of tables, use the `merge` function.

To user guide

See the user guide for a full description of the various [facilities to combine data tables](#).

Previous
◀ [How to reshape the layout of tables?](#)

Next
[How to handle time series data with ease?](#) ▶